

Project 1: IaaS—Amazon Web Services

CSE 546: Cloud Computing

1. Summary

In the first project, we will build an elastic application that can automatically scale out and in on demand and cost-effectively by using cloud resources. Specifically, we will build this application using the cloud resources from Amazon Web Services (AWS). AWS is the most widely used IaaS provider and offers a variety of different compute, storage, and message cloud services. Our application will offer a meaningful cloud service to users, and the technologies and techniques that we learn will be useful for us to build many others in future.

2. Description

Our cloud app will provide an image recognition service to users, by using cloud resources to perform deep learning on images provided by the users. The deep learning model will be provided to you, but you need to build the application that uses this model to provide the service, and meet the following typical requirements for a cloud application:

- 1) The app should take images uploaded via HTTP and perform image recognition on these images using the provided deep learning model. It should also return the recognition result (the top 1 result from the provided model) to the users via HTTP.

For example, the cloud app takes HTTP requests in the form of

`http://[IP Address]/cloudimagerecognition.php?input=[URL of the image]`

For the above request, the response should be *"bobsled"*

If your app's service address is different (because you do not use PHP), you must specify its format in your project README file.

Your app should provide this service on a public IP address. Specify your app's IP address in your project README file.

To facilitate the testing, a standard image dataset is provided to you at:

`http://visa.lab.asu.edu/cifar-10/`

- 2) The deep learning model will be provided to you in an AWS image (ID: ami-07303b67, Name: imageRecognition, Region: us-west-1a). Your application will invoke this model to perform image recognition on the received images.
- 3) The application should be able to handle multiple requests concurrently. It should automatically scale out when the request demand increases, and automatically scale in when the demand drops. Because we have limited resources from the free tier, the application should use no more than 20 instances, and it should queue all the pending requests when it reaches this limit. When there is no request, the application should use the minimum number of instances. Give your instances meaningful names, for example, a web-

tier instance should be named in the fashion of “web-instance1” and an app-tier instance should be named “app-instance1”.

- 4) All the inputs (image names) and outputs (recognition results) should be stored in S3 for persistency. They should be stored in a bucket in the form of (input, output) pairs, e.g.,

(0_cat.png, bobsled)

Specify your bucket name in your project README file.

- 5) The application should handle all the requests as fast as possible, and it should not miss any request. The recognition requests should all be correct.
- 6) For the sake of easy testing, use the resources from only the *us-west-1a* region for all your app's needs.

3. Submission

The submission includes two steps:

- 1) You need to submit your implementation on Blackboard by **March 28th 11:59:59pm**. Your submission should be a single zip file that is named by the full names of your team members. The zip file should contain all your source code and the AWS credentials for running your application. It should also contain a simple README file that lists your group member names and the public IP and S3 bucket name of your app; and a detailed PDF file that describes the design of your application and any additional information that can help the instructor understand your code and run your application. A template will be provide to you for writing the report. Do not include any binary file in your zip file. Only one submission is needed per group.

Failure to follow the above submission instructions will cause penalty to your grade. The submission link will be closed immediately after the deadline.

- 2) You need to give a live demo to the TAs on **March 29th** in class. We will try the code that you submitted to Blackboard. You have only five minutes for the demo, and there will be no second chance if you fail the demo.

4. Policies

- 1) Late submissions will **absolutely not** be graded (unless you have verifiable proof of emergency). It is much better to submit partial work on time and get partial credit for your work than to submit late for no credit.
- 2) Each group needs to **work independently** on this exercise. We encourage high-level discussions among groups to help each other understand the concepts and principles. However, code-level discussion is prohibited and plagiarism will directly lead to failure of this course. We will use anti-plagiarism tools to detect violations of this policy.